

Contents

1	Introduction	2
2	Design Goals	2
2.1	Tiered Learning Resources	3
2.2	Music-Forward Resources	3
3	Ecosystem of Learning Materials	4
3.1	Creative Coding Environment Materials	4
3.2	Web Reference	4
3.3	Learn Articles	5
3.4	Explore Articles	5
3.5	Discourse	6
3.6	Inter-connectivity of Resources	6
4	Example Lesson Plan: MLPRegressor	6
4.1	Classifier Extension	8
5	Common Learning Challenges & Strategies	9
5.1	New Ways of Using Buffers	9
5.1.1	Initial Encounter	9
5.1.2	Holding Analyses	10
5.1.3	Manipulating & Copying Data	10
5.1.4	Why Buffers	11
5.2	Stateful Objects	11
5.3	Fourier Transform & STFT	12
5.4	Advanced Neural Networks	12
5.4.1	Web Resources	12
5.4.2	Building Intuition about How Neural Networks Learn	13
5.4.3	MLP Parameters	13
5.4.4	Visualizing a MLP	13
5.5	Threading	13
5.5.1	SuperCollider	13
5.6	Navigating Human & Machine Assumptions	14
5.6.1	Machine Listening: Pitch	14
5.6.2	Statistics	15
5.6.3	“Know Your Data”	16
5.6.4	Scalars & Distance as Similarity	16
5.6.5	Human vs. Machine Assumptions	17
5.7	De-Myth-ifying Machine Learning	17
5.8	CCE Specific Objects	18
6	Relevance to Contemporary Society	19

1 Introduction

The Fluid Corpus Manipulation Toolkit (FluCoMa) enables techno-fluent musicians to use machine listening and machine learning in their creative practice within the familiar environments of Max, SuperCollider, and Pure Data. Housed at the University of Huddersfield’s Center for Research in New Music (CeReNeM) in the United Kingdom, the project’s primary development period was funded by a five-year grant (2017-2022) from the European Research Council.

While the core of the project produced code packages for all three major computer music programming environments, FluCoMa’s vision and impact are broader. FluCoMa is also a collection of learning resources, code examples, commissioned artworks, musicological articles, interviews, podcasts, a philosophy about interface design for creative coding, a conversation about the future of computer music, a curriculum of machine listening and machine learning topics, a community of users around the world, and more. The materials included here extend this ecosystem to encompass resources for pedagogues that might be used to teach FluCoMa in various settings. While some of the ideas and resources presented below are FluCoMa-specific, many of them are toolkit-agnostic and we hope that they can be used by anyone looking to teach machine listening and machine learning for creative music making.

The extent of the FluCoMa ecosystem supports our belief that “providing the tools” is not enough to achieve FluCoMa’s mission: *enable techno-fluent musicians to use machine listening and machine learning in their creative practices*. To truly enable these artists we also must provide knowledge and inspiration, all in a way that is learnable for our main user base: computer musicians. Topics in machine listening, machine learning, computational thinking, and data science are often not included in the training of electronic musicians and therefore a primary objective of FluCoMa is to build bridges of understanding from the knowledge that comes from computer music training to a degree of fluency with these topics that enables creative music making.

2 Design Goals

The learning materials included in this document have been created through participatory, iterative, and interactive design. The first draft of learning materials were developed based on feedback from composers commissioned to create music with early versions of the Toolkit. Their feedback on what was confusing about the tools, resources, examples, etc. led to a second draft of learning materials which were then used in over thirty workshops around the world with computer musicians from various backgrounds. The feedback from these workshop participants further refined the learning materials and have led to the broad ecosystem of learning resources now found as part of the FluCoMa Toolkit.

2.1 Tiered Learning Resources

Because different learners will desire different degrees of fluency with these topics, we have tried to tier the learning resources accordingly. The most proximal resources (such as environment-native help files) provide a working understanding of what an algorithm does alongside musical examples of how it might be used. Additional resources (such as the internet resources) provide a deeper understanding that might satisfy one’s curiosity and/or build an intuition of what is happening “under the hood,” both of which can enable a more informed manipulation of the Toolkit. When appropriate, we also link to the white paper that describes the algorithm from an engineering perspective, should the learner wish to pursue that amount of technical detail.

In some cases, paths for pursuing additional resources are more extensive and less linear. This is especially true for the more complex tools in FluCoMa, such as neural networks which have multiple dedicated web resources with varying degrees of technical information, any one of which might come after an initial introduction, but when taken all together encompass the degree of fluency we propose for our learners and users. Learners who are eager to manipulate the many neural network hyper-parameters might jump to MLP Parameters, while a user who needs a little more time absorbing how a neural network works might opt for MLP Training. (Also see MLPRegressor, MLPClassifier, and Training-Testing Split.)

Tiered learning resources allow the learner to pursue knowledge as far as they deem appropriate to feed their creative practice in a given moment. Providing the learner what they *need* to know *when* they *need* to know it enables them to stay focused on a creative idea and not become overwhelmed by what could be a very large body of knowledge with a daunting learning curve. This sensitivity to the relationship between creative pursuits and technical knowledge reflects earlier findings of FluCoMa outlined as “Techno-Fluency” and “Divergence.” [1] By offering signposts and links to further resources, the user knows where to keep learning if necessary in the moment, or, in the future if they decide to continue exploring.

2.2 Music-Forward Resources

Because of the specificity of our target learner (a creative coding musician), we have always tried to keep our learning materials and examples musically oriented (as can be seen below). We aim to have the help files and example code make sound in a creative way. When possible, we offer pedagogical examples and thought experiments that will feel familiar and relevant to our learner such as instrument samples, drum hits, MIDI notes, synthesizer settings, measures of frequency and loudness, etc. We hope this strategy will not only explain a tool and its interface, but also provide some copy-and-paste code to get started quickly, and generally get the musical creative juices flowing while a user is engaging in the learning process.

3 Ecosystem of Learning Materials

3.1 Creative Coding Environment Materials

Each creative coding environment (CCE) supported by FluCoMa (Max, SuperCollider, and Pure Data) has a native system for offering reference materials. In Max and Pure Data a “help file” provides annotated examples, while a “reference” offers additional description and detail about parameters. In SuperCollider all of this information is contained in one “help file” document. Despite this difference in interface, we have strived to keep the CCE-based FluCoMa materials similar across all three environments. This is enabled, in part, by the shared documents used to render the reference materials for all three CCEs.

The FluCoMa materials provided natively in the CCE are often the learner’s first engagement with our supporting materials. Therefore, the information provided is intended to provide the learner/user a working understanding of what an object does, how it might be used musically through a sound-making example, and if appropriate, how it interfaces with other FluCoMa objects. This information will hopefully provide a learner some motivation for exploring an object and the amount of knowledge necessary to do so. Each resource in the CCE contains a link to the corresponding web-reference if the learner wishes to pursue a deeper understanding of the object.

The example code provided has been created to be as similar as possible across the three CCEs while keeping idiomatic to each environment. One goal of this is to enable cross-environment communication and knowledge sharing. Users in different CCEs are able to discuss the technical and musical facets of a shared FluCoMa example. It is also possible that this allows referencing help files and example code to a classroom containing a diversity of coding environment users. Lastly, this keeps all three CCEs on an equal status, preventing any potential inference of CCE preference within the FluCoMa Toolkit.

3.2 Web Reference

Every object in FluCoMa (except CCE-specific objects) has a web reference found at learn.flucoma.org. Because the reference materials that appear natively in the CCEs link to the web, the web references are considered to be a secondary resource. The goal of the web references is to offer more detailed descriptions of how an algorithm is working “under the hood”. This may be useful for satisfying curious learners and/or building a better intuition for musical uses of an object. In either case, it could lead to a more informed, and therefore potentially more advanced, use of the object.

Many of the web references have interactive explanations that allow a learner to “use” the algorithm in the browser. Not only can these be used by individual learners, we have also found that these are very useful for explanations in the course of teaching. For example, when creating a KDTree for the first time during a code-along class, using the interactive page helps give learners a visual sense of what is happening (especially if the KDTree is about to be used

with the plotter). The MFCC reference has an interactive explanation that invites a learner (or a teacher during demonstration) to step through a series of interactions that build intuition about MFCCs.

We imagine the web references to be used as solo-learning resources, in parallel with class assignments, as teaching demonstrations, and/or as useful reminders.

3.3 Learn Articles

FluCoMa’s learning website (learn.flucoma.org) also contains many articles about topics that may not fit in a single web reference page. These articles may arise as a tertiary step in a learners path and are likely to be encountered after the CCE materials and web reference.

There are a few varieties of articles found in this category:

- explainers specific to a single FluCoMa object that offer a depth of knowledge about the internal algorithms that would be outside the scope of a web reference page, such as Audio Decomposition using BufNMF.
- knowledge about data science that is useful for using many of the FluCoMa objects, such as Distribution and Histograms and Why Scale? Distance as Similarity
- common workflows using the toolkit, such as Batch Processing with FluCoMa

These articles are not necessarily designed to be consumed in series as part of a sequence of learning (although some could be used this way). Instead, each article is made to be approached by a learner (or guided by a teacher) at a particular point in the learning process and revisited as necessary. The idea for many of these articles arose in direct response to questions asked by workshop participants and therefore are designed to answer or provide context to common questions asked by learners. When designing a curriculum or syllabus, many of these articles would support student learning for different topics in a course.

3.4 Explore Articles


In addition to the web reference and learn articles, the website has many additional materials surrounding artistic uses of the toolkit. These include example artworks, interviews with creative coders, and musicological articles that offer in-depth analysis and example patches of music made with FluCoMa. All of these can be used as context, examples, and inspiration for learners. These can also be used as entry points, as many learners will find the work produced using FluCoMa or the musical ideas expressed in these articles inspiring and motivating.

3.5 Discourse

The international community of FluCoMa users primarily communicates through the Discourse online discussion forum. Any learner (or user) of FluCoMa should be a member of the Discourse as it is an excellent place to converse with like-minded artists. Learners may find the search functionality very useful to see if others have already asked and answered a question they have. The community is very positive and supportive, so it is a safe place to ask all kinds of questions. In addition to the thread for , *Getting Started and Wayfinding*, *Usage Questions*, there are also threads for *Code Sharing*, *Learning Resources*, and *Interesting Links* making it another place for learners to browse for examples, inspiration, and knowledge.

3.6 Inter-connectivity of Resources

As described above, the FluCoMa learning resources are generally tiered to offer learners the degree of detail needed at a given moment to pursue a creative idea. One way in which our tiered approach is executed is cross referencing the different learning resources. The help files in each creative coding environment (CCE) link to their respective web reference, from which a learner can be pointed to many more resources. The web reference, learn articles, and explore articles all cross-link with each other so that a learner reading a web reference might discover an Explore Article about a musician’s use of an object, and from there discover a Learn Article to help them pursue the creative use they just learned about, etc. Different learners will need different degrees of technical specificity, inspiration, and modes of engagement at different times. We hope that setting someone “loose” on the website will enable them to find uses of FluCoMa that are meaningful to them as well as the knowledge needed to support them.



learning. The lesson plan summarized here introducing neural networks with the MLPRegressor is often as a first instruction to FluCoMa.

Watch a Video Tutorial of this Lesson Plan

- MLPRegressor in Max
- MLPRegressor in SuperCollider

The first activity that we often engage learners with is building a neural network that performs regression to control a synthesizer with ten control parameters from a space of only two parameters. This gets learners making sound quickly and uses part of the toolkit that is often quite exciting for newcomers to machine learning (neural networks). This activity takes anywhere from 40-90 minutes depending on the class of learners. Here is a brief outline of the lesson plan:

1. Share a real world example (including watching a performance excerpt) of why someone might want to use a system like this.
2. Using a slides presentation step through how we will be collecting training data and training the neural network, including some intuition about how the training process works.
3. Open up the CCE of choice and demonstrate a completed version (Max, SuperCollider) of the instrument we're about to code.
4. Code the instrument together, as a code-along, starting from a "starter patch" (Max, SuperCollider) that has a few key items already in place:
 - a synthesizer to control
 - a 2D control space to use as input to the neural network
 - a MLPRegressor object with many arguments already specified
5. Let the learners play with the instrument (and augment it in their own way).

The starter patch is important here so that we don't spend too much time doing CCE-specific boiler plate code but instead get right into using FluCoMa. It also ensures that learners have a synthesizer to make sound with right away when the code-along is complete. Because there are many arguments to the MLPRegressor object and each of them can require a fair amount of explanation to use well—and in coordination with each other—we've chosen to provide the arguments to the MLPRegressor programmed into the starter patch. During the lesson we tell the learners that these arguments can be explored further in a future lesson and/or in the learn article.

The extensions of this activity are to:

1. Practice training the neural network.
 - clearing the neural network and retraining
 - training the neural network to a different degree to see if it is more (or less, or differently) musically expressive

- delete the input data and choose new data points to pair with the synthesis parameters
 - delete all the data and create a whole new training
2. Attach a different sound-making algorithm to the output of the neural network
 - granular synthesis / sample playback
 - frequency modulation
 - a VST plugin
 - we have often encouraged learners to bring a sound-making algorithm of their design to the workshop to connect as a next step to this activity
 3. Attach a different type of controller to the input of the neural network. This might be something like:
 - multiple parameters on TouchOSC
 - MIDI controller
 - leap motion
 - wearable device
 - pixel information from a camera (perhaps using Jitter in Max)

Not only does this activity quickly provide learners with a machine learning instrument that is very extensible, it also introduces some of the key elements of FluCoMa:

1. DataSets
2. Buffer interfacing
 - `fluid.buf2list~` and `fluid.list2buf~` for Max
 - `FluidBufToKr` and `FluidKrToBuf` for SuperCollider
 - `array set` and `array get` in Pure Data
3. Using small, personalized, artist-created DataSets
4. Aesthetic evaluation of results
5. Iterative trial-and-error workflows with machine learning algorithms

4.1 Classifier Extension

After completing the MLPRegressor activity, one common extension is to do an activity training the MLPClassifier to distinguish between two timbres. Depending on the group of learners and how much time is available, sometimes this would only include opening up the classifier demonstration file (Max, SuperCollider) and doing a quick training and testing, along the way relating it to what was just done with the MLPRegressor activity. If time allowed and there is interest we performed a more involved activity similar to the MLPRegressor: demoing the code and then building it together.

Watch a Video Tutorial of this Lesson Plan

- MLPClassifier in SuperCollider
 - MLPClassifier in Max
1. Share a real world example (including watching a performance excerpt, this one has a before and after) of why someone might want to use a system like this.
 2. Using a slides presentation step through how we will be collecting training data and training the neural network, including some intuition about how the training process works.
 3. Open up the CCE of choice and demonstrate a completed version (Max, SuperCollider) of what we're about to code.
 4. Code-along, starting from a “starter patch” (Max, SuperCollider) that has a few key items already in place:
 - the sound files containing timbres that will be used as training and testing data
 - a MLPClassifier object with many arguments already specified (for the same reason described above)
 5. Let the learners explore classifying some of their own sounds

5 Common Learning Challenges & Strategies

Over the course of teaching many workshops, we observed some common challenges for FluCoMa learners. Below are a few of the challenges we found and some strategies for approaching them pedagogically.

5.1 New Ways of Using Buffers

FluCoMa uses buffers to store all kinds of data, not just audio. This may be new for learners who are used to using buffers *only* for holding audio, and may even associate the two as a single concept (“buffer == audio”). This becomes increasingly complicated when we begin to manipulate the data in buffers as arrays and matrices.

5.1.1 Initial Encounter

The first moment at which a learner is asked to view a buffer in a new way is often when we allocate a buffer to hold a data point. If the data point has just 2 dimensions (such as with the MLPRegressor activity), we will allocate the buffer with only 2 frames (in Max: `@samps 2`; in SuperCollider: `Buffer.alloc(s,2)`). At this point we try to offer something like the following:

“The way we most often use buffers in [this CCE] is to hold audio samples which very commonly have 44,100 samples per second, so our buffers could have many tens or hundreds of thousands of values in them. Because I know this buffer is only going to need to hold two values, I’ll allocate it to have just two samples.”

5.1.2 Holding Analyses

Once we start writing audio analyses into buffers (with the **feature** argument), learners often have a hard time keep track of the structure of the buffers (What do the channels represent? How many are there? What do the frames represent? How many are there?). We found that offering CCE-agnostic charts of the “shape” of the buffer is very helpful for giving learners a mental model to hold in their mind.

FluidBufStats writes the analysis to *another* buffer

frame:		0	1	2	3	4	5	6
analysis feature ->	chan: 0	mean of chan 0	stand. dev. of chan 0	skewness of chan 0	kurtosis of chan 0	low (min) of chan 0	mid (median) of chan 0	high (max) of chan 0
	1	mean of chan 1	stand. dev. of chan 1	skewness of chan 1	kurtosis of chan 1	low (min) of chan 1	mid (median) of chan 1	high (max) of chan 1
	⋮

Figure 2: Chart demonstrating what the channels and frames of a particular buffer represent.

It’s also useful to point out that for buffers that hold audio analyses, the frames (or what we sometimes refer to as the “x axis” in reference to the charts above) is still a time series, just like audio is, but now it’s not a time series of voltages (as in audio), it’s a time series of descriptors (such as spectral centroid).

Because each frame represents an FFT frame from the STFT analysis, the sample rate would not be a usual 44,100 samples per second, but a much lower rate of frames per second (FFT frames per second). For example, if an audio buffer with a sample rate is 44,100 Hz is analyzed with a **hopSize** of 512 samples, the **features** buffer that the analyses get written into will have a sample rate of 86.1328125 frames per second ($44100 / 512$). In SuperCollider and Max, FluCoMa buffer processors (such as the audio descriptor analyzers) set the *sample rate* of these buffers appropriately (Pure Data arrays don’t hold metadata). If the values in that buffer are read back at that rate, they will correspond in time (be synchronized with) to the audio on which the analysis was based. Point this out to learners helps them remember and conceptualize the relationship between the source audio being analyzed, the STFT process, and the resulting time series of descriptor values.

5.1.3 Manipulating & Copying Data

Often it is necessary to manipulate the data in a buffer, such as pick out values from certain channels and/or frames and copy them to another buffer. In

order to provide some “test and check” interactivity to build fluency with these operations, the appropriate web references have interactive GUIs for practicing:

- BufSelect
- BufSelectEvery
- BufFlatten
- BufCompose
- BufScale

5.1.4 Why Buffers

It also may be of interest for learners to hear the explanation of *why* buffers are used in this way. FluCoMa uses buffers in this way for a few reasons:

- The notion of “buffer” is shared across all three CCEs that FluCoMa supports. This allows for shared syntax and usage of objects across all three environments.
- In all three CCEs, buffers are accessed at the lower levels of code allowing for:
 - Faster processing by interfacing directly with the C++ code.
 - Simpler implementation of functions across all three CCEs because all three environments share the same C++ code base for all of the audio analysis, buffer processing, and algorithms.
- Having data in buffers allows for it to be more flexibly accessed and used in other parts of the CCE. For example because the MLPRegressor writes predictions into a buffer, it’s possible to be predicting wavetable shapes directly into a buffer that is simultaneously being read out of, all on the scsynth server.

5.2 Stateful Objects

Many of the FluCoMa data objects hold some state. For example after calling `fit` (or `fitTransform`) on a Normalize object, it holds the minimum and maximum value of each dimension in the fitted DataSet so that it can scale future `transform` calls appropriately. MLP objects hold the state of the MLP internal parameters (the state of the trained model). This interface design is based on many of the data processing objects in the Python sci-kit learn package. We found that some FluCoMa learners find it challenging to conceptualize or remember that certain objects are holding a state that they will need to call upon later. One feature that may help with conceptualizing objects in this way is the option to *name* objects in Max and Pure Data (SuperCollider natively uses variable names to identify objects).

Named objects may help learners remember that certain objects hold state because they have a sense of it being a non-generic, task-specific object, such as a Normalize object called `norm-pre-pca`. This gives it a special sense of purpose and an indicator of what state it holds and where in data processing one would call upon that state.

5.3 Fourier Transform & STFT

As with many audio tasks, the STFT is central to much of how a learner interacts with FluComa. Many times learners can do exciting things, learn a lot about the toolkit, and make great music without reflecting on the FFT processes happening “under the hood.” We have found however, that for many of the algorithms in FluCoMa (such as AudioTransport, Sines, and many more), adjusting STFT settings (`windowSize`, `hopSize`, and `fftSize`) has an important aesthetic impact on the results, and therefore we suggest that it is important to understand what impact these parameters have.

There exists a wealth of resources on the internet for learners to build fluency with the Fourier Transform, so we didn’t feel the need to recreate many of these learning tools. We have however, curated a small set of musician-oriented ideas that will be important for FluCoMa learners to consider as they are approaching the toolkit, which exist on these two pages:

- Fourier Transform
- BufSTFT

5.4 Advanced Neural Networks

Learners often follow up the MLPregressor Lesson Plan with questions about the hyper-parameters (which FluCoMa calls parameters or arguments) of the MLP. In shorter workshops (two days or fewer) we have felt that it’s not enough time to delve into this with enough depth to make it well understood and *useable* for the participants, so we’ve directed them towards our web resources on the topic. In longer workshops (three days or more) we have taken time later in the week (day three or four) after the introductory activity to unpack many more ideas and strategies about the MLP objects.

5.4.1 Web Resources

- Neural Network Training is an overview of how neural networks learn. It is intended for those who would benefit from gaining a little more intuition about what is going on “under the hood” or for learners that have a little more curiosity they want to satisfy. Much of what this article expresses is included in the introductory activity.
- Neural Network Parameters goes through each parameter in the MLP objects and gives a more thorough description of what it controls, why one might adjust it, and what a generally reasonable starting place is. This is often where we direct learners who ask about these parameters when we don’t have time to unpack them during a workshop.
- Training and Testing Data describes why it can be important to validate the results of a trained MLP. It explains why one would go about validating a model, what to look out for, what certain results might mean, and what one might do to improve a model.

5.4.2 Building Intuition about How Neural Networks Learn

The sequence of explanation that we’ve used for both the MLPRegressor and MLPClassifier seems to work quite well for giving learners intuition about the training process of an MLP. These explanations can be seen at the beginning of the the MLPRegressor and MLPClassifier tutorial videos.

5.4.3 MLP Parameters

As stated above, when we have had time in workshops, we’ve allocated time to explain the parameters in more detail. A very useful site for explaining and playing with the `momentum` parameter can be found on distill.

5.4.4 Visualizing a MLP

One might notice in many of the resources above that there are node-and-edge graphs of MLP architectures. We found that these are very useful for learners to visualize and concretize a few facets of MLPs: (1) “feed-forward”, (2) “back-propagation”, (3) “fully-connected layers”, (4) numbers of hidden layers and nodes, (5) total number of parameters in an architecture, and more. We also learned that it is important to have a visual representation of the architecture that is *actually used* in the activity. We use a very basic graphviz script to generate these graphs.

5.5 Threading

5.5.1 SuperCollider

In SuperCollider, many FluCoMa workflows rely heavily on the server’s OSC queue. It will be important for FluCoMa learners to become familiar with how the server processes the OSC messages it receives.

In particular, when performing many buffer operations, a common workflow is to use the `processBlocking` call for all of the operations which will place each operation in the server’s OSC queue. This allows for very fast buffer processing because the the operations are all lined up in order and the server moves on to the next one as soon as it completes with the previous (all without any unnecessary `sync` with the language). One way in which this often confuses learners is that if these operations are inside of a loop (that is iterating over slice points perhaps), and the code is printing *which* slice is currently being analyzed (see below), it will report that it has “completed analyzing” all the slices, when in fact all this suggests is that all of the `processBlocking` operations have been sent to the server.

```
fa.doAdjacentPairs{
  arg start, end, i;
  var num = end - start;

  start.postln;
```

```

end.postln;
i.postln;

// analyze a sound slice for spectral centroid
FluidBufSpectralShape.processBlocking(s,~src,start,num,features:spec,select:[\centroid]);

// get the mean centroid for this sounds slice
FluidBufStats.processBlocking(s,spec,stats:stats,select:[\mean]);

// copy the mean spectral centroid to the appropriate index (destStartFrame) of
// the buffer called meancentroids
FluidBufCompose.processBlocking(s,stats,destination:meancentroids,destStartFrame:i);

"completed analyzing slice %".format(i).postln;
};

```

One good way of clarifying this for learners is to enforce a **sync** between the language and the server every 100 slices or so using a line of code like `if((i % 100) == 99){ s.sync }`. This will allow the code to post “completed” for every 100 slices, but wait until those slices are completed before posting anything else. This clarifies the speed at which the analyses are actually happening, while adding a trivial amount of processing time (due to syncing) to the overall analysis.

One could sync each time the loop is executed (on every slice) but this would significantly slow down the analysis.

5.6 Navigating Human & Machine Assumptions

One of the most challenging conceptual hurdles for FluCoMa learners is to reconcile the differences between the way machines and humans listen. What perceptually might seem obvious to a human listener can be very challenging for a machine to discern. Newcomers to machine listening often set out to perform a task making many assumptions about how a system will work, what data they will use, and how they will compute a result, not realizing that what they *think* they’re telling the machine or asking it to compute is quite different from what it will give back in return.

5.6.1 Machine Listening: Pitch

One activity that has been quite successful is a simple “listening test.” We ask a class of learners to sing the pitch of a sound file. Most listeners will sing the right pitch class of the tone but down a few octaves from the actual frequency. The first thing to point out is that they were only singing the pitch from the part of the sound file that was *most* pitched. They didn’t even attempt to sing the “pitch” during the scratchy parts. We then look at a chart showing the result of a Pitch analysis on the buffer. One can see where the pitch is stable (the parts

that the listeners sung), but also that there is a lot more “pitch” analysis there. The machine listens to all of it (and reports back on all of it). This is because as humans we’re constantly engaged in multi-modal listening, switching between different ways of perceiving sound, depending on which seems appropriate or useful at a given moment. When the sound file is making scratchy sounds, we as humans don’t even register it as a “pitch” to sing, but a machine does.




Figure 3: Pitch analysis of a buffer that has some pure tone and some scratchy timbres.

Another useful outcome of this activity can be acknowledged when discussing Distance as Similarity. Most listeners will sing the correct pitch class but down a few octaves. For these listeners, being 12 half steps lower (a distance of 12) is closer than being 1 half step lower (a distance of one). This is again a recognition that what humans might assume to be *similar* a machine might not (in this case using Chroma analysis may help align human and machine listening).

5.6.2 Statistics

Many workflows in FluCoMa require the use of statical summaries of audio descriptor time series, real time audio analyses, or whole DataSets. It is important for learners to develop a sense of what tools are available and why one might reach for one statistical summary rather than another.

BufStats is perhaps the most commonly used object in FluCoMa and therefore has a somewhat involved reference page. Many of the statistics available might be familiar to learners (mean, median, minimum, and maximum), while others might be new (standard deviation, skewness, kurtosis, derivatives, etc.). During the introductory tutorials, many moments arise that are useful for reflecting on the statistical analyses being used and how they affect the sonic results being heard. A few examples are reflecting on what it means to have an average spectral centroid of a sound slice and using the maximum value of an analysis rather than the mean.

BufStats has many more features that are somewhat less explored and probably not appropriate for learners just getting acquainted with FluCoMa. There are few Learn Articles that cover these topics in musicianly ways including Weighting Stats and Outliers. It is also important to convey to learners, as is stated on the BufStats page,

While it can be difficult to discern how to use some of these anal-

yses practically (i.e., what does the kurtosis of the first derivative of spectral centroid indicate musically?), these statistical summaries can sometimes represent differences between analyses that dimensionality reduction and machine learning algorithms can pick up on. Including these statistical descriptions in training or analysis may provide better distinction between data points.

5.6.3 “Know Your Data”

As with all data science and machine learning, understanding what data represents, what it can tell you (and more importantly what it *can't*), and what transformations *do* to data is essential. It is continuously important for FluCoMa learners to reflect on their data. There are a few visualization tools that are very important for users to get comfortable with including the Waveform object (`fluid.waveform~` in Max and `FluidWaveform` in SuperCollider). While teaching, these tools should be used whenever possible to help learners understand the data processing that is happening and get them in the habit of visually checking on their data regularly to build understanding of what their data represents.

Many machine learning algorithms make various assumptions about data (similar to how humans make assumptions about sound and machine listening). One of these assumptions is that the dimensions in a `DataSet` are identically distributed, often Gaussian distributed. It can be important to know how one's data is distributed and it is possible to check on using a histogram. Our [Distribution and Histograms](#) page gives some examples of different kinds of distributions, what they mean, and some example code to check on a distribution using a histogram.

5.6.4 Scalers & Distance as Similarity

Another important concept for learners to understand is how measures of distance impact perceptions and assumptions about similarity. Once the machine has “listened” and a statistical summary has been computed, a next step is often to compare data points by computing the distance between them (most of the FluCoMa tools use Euclidean Distance). Computing distance very quickly makes questions about ranges and scaling relevant, such as how a mismatch of scale may overly weight the importance of dimensions that have larger ranges.

This is a great opportunity to compare scalers available in FluCoMa. One way of clearly demonstrating that different scalers will have different sonic results (and that those sonic results are not always *predictable*) is to choose a single point in a `DataSet` (such as one sound slice), and find what the nearest neighbor is with (1) no scaling, (2) `Normalize`, (3) `Standardize`, and (4) `RobustScale`. (This is essentially what the sequence of images does in the [Comparing Scalers](#) page.) Doing these comparisons in real time while hearing the sonic differences can help concretize the importance for learnings.

It can often be important for learners to keep track of which dimensions

might be logarithmic or linear and know how those differences could affect measures of distance. One concrete example we provide is on our scaling page under the heading Linear vs Logarithmic Scales where we state:

For example, frequency analyses might be provided in hertz (which is on a linear scale), however this doesn't reflect how humans actually perceive pitch distance. For a more perceptually relevant scale it is displayed logarithmically in pitch space (perhaps labeled as MIDI notes or semitones). If measuring in hertz, the distance from the A4 down one octave to A3 is 220 hertz, while the distance from A4 up one octave to A5 is 440 hertz—twice as far even though we perceive them to both be one octave! Measuring these distances in semitones however will reflect the way we perceive them: both a distance of 12.

FluCoMa objects which analyze for frequency have an argument called `unit` which specifies if the frequency estimation should be returned in Hz or MIDI notes.

5.6.5 Human vs. Machine Assumptions

One more concrete example of how human and machine assumptions differ comes from a learner who has having trouble getting KMeans to cluster data points in the way they thought it should. The learner wanted KMeans to cluster points on a 2D plot according to the clusters that are easily visually identifiable by a human. KMeans was clustering it much differently, including leaving many clusters empty. This was solved by demonstrating (including with the original data) how the learner could seed KMeans as a “human in the loop” approach to direct its processing with the information also inferred by a human.

5.7 De-Myth-ifying Machine Learning

Sometimes we have questions from learners that sound something like, “I want X to do Y. How can FluCoMa do this?”. At this point, our pedagogical step is to break down the goal into smaller and more specific questions and tasks that we can start approaching together with the learner. This process often reveals the assumptions that the learner might be making about how audio analyses work, or what a machine will be able to perceive, or how long an analysis or algorithm might take, and all of the tradeoffs involved in making decisions about the process. Sometimes the question transforms from “How can FluCoMa do this?” to “Can FluCoMa do this?” at which point perhaps there’s a different tool that we can point them to, or help them realize that their goal is too lofty—that it stems from a belief that “AI can do anything” or “throw it at a neural network and it’ll figure it out”. Usually this process enriches the learner’s ideas about what is possible with FluCoMa (even if it’s not necessarily what they hoped) and provides a lot of possibilities for investigation.

Another de-myth-ification that has occurred is when learners will assume that the machine learning *is* performing some magic. This is often in the form of learners not *validating* or testing the machine learning model or the results of their algorithm. The first disclaimer to make is that, as artists, we're interested in artistically compelling experiences, so regardless of what the algorithm is or isn't doing, if it sounds good, keep it. It is also important to test the systems that we build and see if they are doing what we think they're doing. This is important for a few reasons:

- Validation can reveal our assumptions and/or misunderstandings about *how* things work, providing opportunities to deepen our knowledge and skills.
- Validation can offer ways to improve our system to get even closer to our desired outcome.
- Validation can reveal nuances in the system that might offer more paths of exploration and creativity.

Framing validation with these benefits in mind can help encourage learners to put in the extra work that it takes.

5.8 CCE Specific Objects

Pedagogues should be aware that there are a few objects in FluCoMa that are CCE specific in name and/or implementation because of CCE differences. These objects diverge from FluCoMa's philosophy of cross-environment parity in order to ensure workflows in the toolkit are idiomatic and *Fluid* for beginner and expert users. When teaching a room that has learners using multiple CCEs it is important to be highlighting the differences in the course of the lesson to ensure understanding. The main differences regard interfacing with Buffers.

Writing control information into a buffer

Max	fluid.list2buf
SuperCollider	FluidKrToBuf
Pure Data	(native)

Reading control information out of a buffer

Max	fluid.buf2list
SuperCollider	FluidBufToKr
Pure Data	(native)

The FluCoMa Plotter object has similar functionality and syntax across all three CCEs but divergent implementation

Max	fluid.plotter
-----	---------------

SuperCollider	FluidPlotter
Pure Data	fluid.plotter

The FluCoMa Waveform object has similar functionality and syntax across all three CCEs but divergent implementation

Max	fluid.waveform~
SuperCollider	FluidWaveform
Pure Data	fluid.waveform

6 Relevance to Contemporary Society

We believe that learning about data, data science, and machine learning through FluCoMa can be used as a lens to consider how these tools operate in contemporary society, in particular to uphold inequalities, injustices, and hegemonies. By gaining fluency and understanding with these algorithms, one can come to understand what these algorithms are good for, what they are not good for, how they go wrong, and the relationship between data, algorithms, and the humans that use them. The skills and knowledge, both explicit and tacit, that working with FluCoMa fosters can be used to reflect on many of the AI events and concerns that are constantly appearing in the news. Pedagogues might draw on these events to use as discussion topics where a classroom of learners can collectively reflect on contemporary topics using the experience and understand built through FluCoMa.

Below is a list of books (in no particular order) that provide many examples of contemporary technologies negatively impacting marginalized communities. Many additionally offer directions for how to approach and use data science ethically. We recommend selecting a book or selected readings from these books to augment learning. Each of these is written for different audiences, so selecting which is best is at the instructor/learner's discretion.

- *Data Feminism* by Catherine D'Ignazio and Lauren F. Klein.
- *Weapons of Math Destruction* by Cathy O'Neil
- *Hello World* by Hannah Fry
- *Revolutionary Mathematics* by Justin Joque
- *Blockchain Chicken Farm: And Other Stories of Tech in China's Countryside* by Xiaowei Wang
- *The Alignment Problem* by Brian Christian

References

- [1] Owen Green, Pierre Alexandre Tremblay, and Gerard Roma. Interdisciplinary research as musical experimentation: A case study in musicianly approaches to sound corpora. In *Electroacoustic Music Studies Network*

Conference: Electroacoustic Music: Is it Still a Form of Experimental Music? Zenodo, 2019.